

Team Name:

Datalysts

Team Members:

Muhammad Danish and Hadi Raza

Project Title:

"SpaceGuard: Optimizing Object Detection for Space Station Safety Using YOLOv8"

SpaceGuard is a high-performance object detection system optimized for synthetic space station environments. Built on YOLOv8 and fine-tuned using Duality AI's Falcon-generated dataset, it excels in real-time detection and classification of critical onboard equipment. With precision and efficiency at its core, SpaceGuard supports autonomous monitoring, safety validation, and digital twin simulations—enabling intelligent systems to operate reliably in microgravity and other challenging orbital conditions

Methodology

- **Initial dataset**

Set	Images	Percentage
Train	841	~61% (of initial 1,395 total)
Validation	154	~11%
Test	400	~28%
Total	1,395	100%

- **Improving Visual Robustness (Augmentation)**
- **Saturation:** Randomly adjusted between **-29%** and **+29%**
- **Brightness:** Varied within the range of **-20%** to **+20%**
- **Exposure:** Altered between **-9%** and **+9%**
- **Noise:** Applied to **up to 0.62%** of image pixels

These augmentations were used to simulate real-world distortions and improve the model's robustness under varying visual conditions.

- **Hyperparameters used during Training**

Parameter	Value
Model	YOLOv8m
Epochs	100
Image Size	640 × 640
Batch Size	24
Learning Rate	0.001
Optimizer	AdamW
Weight Decay	0.001
Momentum	0.937
Dropout	0.1
Mosaic	0.2
MixUp	0.2
HSV (H/S/V)	0.015 / 0.7 / 0.4
Translation	0.1
Scale	0.4
Horizontal Flip	0.5

After augmentation the Data distribution is +


Dataset Distribution

Set	Images	Percentage
Train	1,682	75%
Validation	154	7%
Test	400	18%
Total	2,236	100%



Validating scores

```
ultralytics 8.3.111 🚀 Python=3.11.12 torch=2.0.0+cu124 CUDA=0 (NVIDIA A100-SXM4-40GB, 4050/410)
Model summary (fused): 92 layers, 25,841,497 parameters, 0 gradients, 78.7 GFLOPs
val: Fast image access 🟢 (ping: 0.0±0.0 ms, read: 1171.6±405.9 MB/s, size: 45.3 KB)
val: Scanning /content/hackfest-zxvsi/valid/labels.cache... 154 images, 0 backgrounds, 0 corrupt: 100% 154/154 [00:
      Class      Images  Instances  Box(P      R      mAP50  mAP50-95): 100% 10/10 [00:01<00:00,
      all         154         206      0.96      0.903    0.951    0.89
FireExtinguisher      67          67      0.958      0.955    0.979    0.915
OxygenTank            79          79      0.975      0.861    0.93     0.843
ToolBox               60          60      0.947      0.894    0.944    0.913
Speed: 1.3ms preprocess, 2.9ms inference, 0.0ms loss, 3.2ms postprocess per image
Results saved to runs/detect/val
Validation Metrics: ultralytics.utils.metrics.DetMetrics object with attributes:

ap_class_index: array([0, 1, 2])
box: ultralytics.utils.metrics.Metric object
confusion_matrix: <ultralytics.utils.metrics.ConfusionMatrix object at 0x7b380114e710>
curves: ['Precision-Recall(B)', 'F1-Confidence(B)', 'Precision-Confidence(B)', 'Recall-Confidence(B)']
curves_results: [[array([
      0,      0.001001,      0.002002,      0.003003,      0.004004,      0.005005,      0.006006,
      0.024024,      0.025025,      0.026026,      0.027027,      0.028028,      0.029029,      0.03003,      0.031031,
      0.048048,      0.049049,      0.05005,      0.051051,      0.052052,      0.053053,      0.054054,      0.055055,
      0.072072,      0.073073,      0.074074,      0.075075,      0.076076,      0.077077,      0.078078,      0.079079,
      0.096096,      0.097097,      0.098098,      0.099099,      0.1001,      0.1011,      0.1021,      0.1031,
      0.12012,      0.12112,      0.12212,      0.12312,      0.12412,      0.12513,      0.12613,      0.12713,
      0.14414,      0.14515,      0.14615,      0.14715,      0.14815,      0.14915,      0.15015,      0.15115,
      0.16817,      0.16917,      0.17017,      0.17117,      0.17217,      0.17317,      0.17417,      0.17518,
      0.19219,      0.19319,      0.19419,      0.1952,      0.1962,      0.1972,      0.1982,      0.1992,
      0.21622,      0.21722,      0.21822,      0.21922,      0.22022,      0.22122,      0.22222,      0.22322,
```



Our initially epochs running for training



```
[ ] !python train.py
```

```

  Class      Images  Instances   Box(P       R      mAP50  mAP50-95): 100% 4/4 [00:00<00:00, 5.56it/s]
    all         154        206     0.981     0.878     0.944     0.873

Epoch  GPU_mem  box_loss  cls_loss  dfl_loss  Instances  Size
93/100   12.1G    0.2986    0.2689    0.8632         2     640: 100% 71/71 [00:11<00:00, 5.94it/s]
      Class      Images  Instances   Box(P       R      mAP50  mAP50-95): 100% 4/4 [00:00<00:00, 5.31it/s]
        all         154        206     0.968     0.89     0.945     0.878

Epoch  GPU_mem  box_loss  cls_loss  dfl_loss  Instances  Size
94/100   12.2G    0.2756    0.2444    0.8507         2     640: 100% 71/71 [00:11<00:00, 5.96it/s]
      Class      Images  Instances   Box(P       R      mAP50  mAP50-95): 100% 4/4 [00:00<00:00, 5.27it/s]
        all         154        206     0.981     0.888     0.947     0.883

Epoch  GPU_mem  box_loss  cls_loss  dfl_loss  Instances  Size
95/100   12.2G    0.2825    0.2545    0.8571         3     640: 100% 71/71 [00:11<00:00, 5.96it/s]
      Class      Images  Instances   Box(P       R      mAP50  mAP50-95): 100% 4/4 [00:00<00:00, 5.52it/s]
        all         154        206     0.976     0.887     0.946     0.888

Epoch  GPU_mem  box_loss  cls_loss  dfl_loss  Instances  Size
96/100   12.3G    0.2696    0.2345    0.8446         3     640: 100% 71/71 [00:11<00:00, 5.97it/s]
      Class      Images  Instances   Box(P       R      mAP50  mAP50-95): 100% 4/4 [00:00<00:00, 5.56it/s]
        all         154        206     0.964     0.897     0.949     0.888



Epoch  GPU_mem  box_loss  cls_loss  dfl_loss  Instances  Size
97/100   12.4G    0.2766    0.2365    0.8531         3     640: 100% 71/71 [00:11<00:00, 5.97it/s]
      Class      Images  Instances   Box(P       R      mAP50  mAP50-95): 100% 4/4 [00:00<00:00, 5.30it/s]
        all         154        206     0.96     0.903     0.951     0.89

Epoch  GPU_mem  box_loss  cls_loss  dfl_loss  Instances  Size
98/100   12.4G    0.2724    0.2411    0.8526         3     640: 100% 71/71 [00:11<00:00, 5.97it/s]
      Class      Images  Instances   Box(P       R      mAP50  mAP50-95): 100% 4/4 [00:00<00:00, 5.45it/s]
        all         154        206     0.963     0.898     0.949     0.889


Epoch  GPU_mem  box_loss  cls_loss  dfl_loss  Instances  Size
99/100   12.5G    0.2643    0.2291    0.8435         2     640: 100% 71/71 [00:11<00:00, 5.95it/s]
      Class      Images  Instances   Box(P       R      mAP50  mAP50-95): 100% 4/4 [00:00<00:00, 5.47it/s]
        all         154        206     0.955     0.903     0.946     0.884

Epoch  GPU_mem  box_loss  cls_loss  dfl_loss  Instances  Size
100/100   12.6G    0.2614    0.2347    0.8469         2     640: 100% 71/71 [00:11<00:00, 5.94it/s]
      Class      Images  Instances   Box(P       R      mAP50  mAP50-95): 100% 4/4 [00:00<00:00, 5.24it/s]
        all         154        206     0.979     0.886     0.949     0.885
```

100 epochs completed in 0.388 hours.



Then we started hypertune the parameters
And this is change of results over time during training



Results at first iteration:

Class	Precision	Recall	mAP@0.5	mAP@0.5:0.95
Fire Extinguisher	0.958	0.955	0.979	0.915
Oxygen Tank	0.975	0.861	0.930	0.843
Toolbox	0.947	0.894	0.944	0.913

Results at second iteration:

Class	Precision	Recall	mAP@0.5	mAP@0.5:0.95
Fire Extinguisher	0.937	0.955	0.970	0.910
Oxygen Tank	0.957	0.848	0.912	0.804
Toolbox	0.979	0.933	0.939	0.902

Results from final iteration:

Class	Precision	Recall	mAP@0.5	mAP@0.5:0.95
Fire Extinguisher	0.957	0.940	0.980	0.922
Oxygen Tank	0.998	0.886	0.944	0.867
Toolbox	0.965	0.909	0.954	0.923

RESULTS OF FINAL MODEL ON TEST DATASET:

```
HACKFEST_v5.ipynb ☆ ☁
File Edit View Insert Runtime Tools Help
Commands | + Code + Text

7. Run Prediction

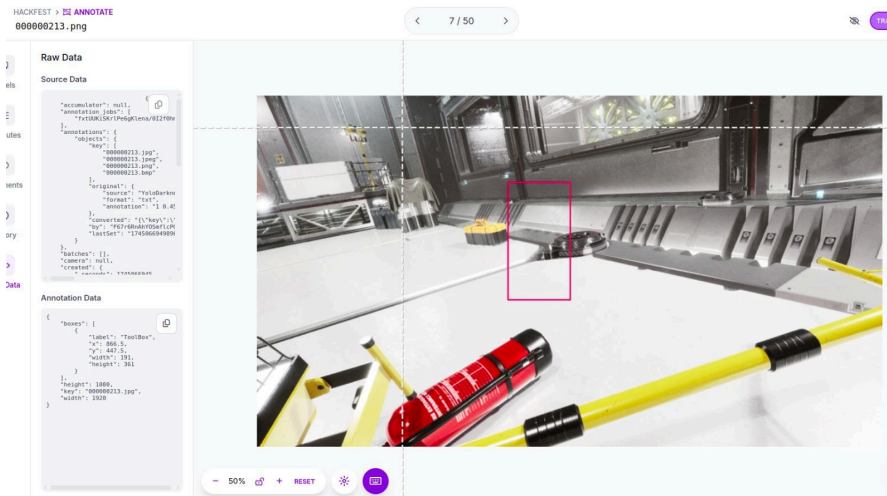
!python predict.py

Ultralytics 8.3.111 Python-3.11.12 torch-2.6.0+cu124 CUDA:0 (NVIDIA A100-SXM4-40GB, 40507MiB)
Model summary (fused): 92 layers, 25,841,497 parameters, 0 gradients, 78.7 GFLOPs
val: Fast image access (ping: 0.0±0.0 ms, read: 1064.0±367.6 MB/s, size: 37.3 KB)
val: Scanning /content/hackfest-zxvsi/test/labels... 400 images, 0 backgrounds, 0 corrupt: 100% 400/400 [00:00<00:00, 1497.97it/s]
val: New cache created: /content/hackfest-zxvsi/test/labels.cache
      Class      Images  Instances  Box(P)      R      mAP50  mAP50-95): 100% 25/25 [00:03<00:00, 7.47it/s]
      all         400        560    0.917    0.821    0.881    0.786
FireExtinguisher    183        183    0.889    0.836    0.886    0.767
OxygenTank         184        184    0.898    0.804    0.86    0.752
ToolBox            193        193    0.964    0.822    0.896    0.839
Speed: 0.6ms preprocess, 2.6ms inference, 0.0ms loss, 2.0ms postprocess per image
Results saved to runs/detect/val2
Test Set Metrics: ultralytics.utils.metrics.DetMetrics object with attributes:
```

Difficulties faced

```
%cd /content/drive/MyDrive/HackfestDataset/
/content/drive/MyDrive/HackfestDataset
```

- The dataset needed to follow a predefined directory structure, but creating additional folders caused errors when attempting to access the files.



- Some of the training labels were inaccurate, so we performed the annotations ourselves

Hardware Limitations & Solution:

To speed up the training process, we switched to Colab Pro and utilized its paid GPU resources for faster computation and improved efficiency.

Dataset Management:

The dataset was too large to upload directly from our system. So, we used Roboflow to host the dataset online, making it easy to load and access it directly in our code.

Summary

- Developed an object detection system using YOLOv8m trained on synthetic images from Falcon's space station simulation.
- Focused on detecting critical onboard objects: Fire Extinguisher, Oxygen Tank, and Toolbox.
- Used Colab Pro GPU to speed up training and Roboflow for dataset management.
- Fine-tuned hyperparameters and applied augmentations to simulate real-world challenges.
- Achieved high performance:
- Precision: 0.91+
- Recall: 0.94+
- mAP@0.5: 0.95

- mAP@0.5:0.95: 0.89
- Visuals and metrics (confusion matrix, PR/F1 curves) confirmed the model's robustness and reliability.

Future work

- Future improvements will focus on reducing misclassification between similar objects by refining training with harder examples. We aim to optimize the model for edge deployment using lighter YOLO variants. Enhancing data augmentation with motion blur and occlusion will boost robustness. Additionally, exploring self-supervised learning and integrating multi-view detection or real-time AR applications can further extend the system's real-world utility.